

# High-Level Design (HLD) – PDF RAG System

Project Title: Retrieval-Augmented Generation (RAG) System for PDF Documents

Author: Anurag Mishra

Date: 2025-09-01

## 1. Project Overview

The goal of this project is to develop an end-to-end RAG system that can ingest PDF documents (native + scanned), perform retrieval-based question answering, generate answers with citations, and highlight exact evidence in the source pages. The system is fully local, using open-source tools like SentenceTransformers, ChromaDB, PyMuPDF, and LLaMA3.

## 2. System Architecture

Components:

### 1. PDF Ingestion

- Input: Raw PDFs (data/raw\_pdfs/)
- Output: pages.jsonl (text and metadata per page)
- Handles both native PDFs and scanned PDFs with OCR

### 2. Chunking

- Splits text into manageable chunks (default: 1000 chars with 200 char overlap)
- Output: chunks.jsonl with chunk IDs, page numbers, pdf\_path

### 3. Embedding + Indexing

- SentenceTransformer (all-MiniLM-L6-v2) generates embeddings per chunk
- Embeddings stored in ChromaDB for fast similarity search

### 4. Retriever

- Query embeddings compared against ChromaDB
- Top-k most relevant chunks returned

### 5. LLM Query

- Prompts constructed with retrieved chunks
- LLaMA3 generates answer with citations [p.N]

### 6. Highlighting

- Evidence snippets highlighted in original PDF using PyMuPDF
- Output: Annotated PDF

## 7. Web UI (Optional)

- Streamlit interface for selecting PDFs, asking queries, and downloading highlighted PDFs

## 3. Module Descriptions

Module	Responsibility	Input/Output
----- ----- -----		
ingest.py	Extracts text + metadata from PDFs	Input: PDF, Output: pages.jsonl
chunk.py	Splits pages into chunks	Input: pages.jsonl, Output: chunks.jsonl
embed_index.py	Generates embeddings and builds ChromaDB index	Input: chunks.jsonl, Output: ChromaDB
rag_pipeline.py	Handles query, retrieval, LLM prompt, and evidence list	Input: query, Output: answer + evidences + annotated PDF
highlight.py	Highlights evidence text in PDF using PyMuPDF	Input: PDF + evidences, Output: annotated PDF
app.py	Streamlit UI for query and download	Input: user query, Output: UI + highlighted PDF

## 4. Data Flow

1. Ingest PDF -> extract pages -> store as JSONL
2. Chunking -> split pages into overlapping chunks -> store as JSONL
3. Embedding & Indexing -> encode chunks -> store in ChromaDB
4. Query -> user asks a question -> embed query -> retrieve top-k chunks -> construct prompt
5. LLM Answer Generation -> LLaMA3 generates answer using prompt
6. Highlight Evidence -> locate evidence in original PDF -> produce highlighted PDF
7. UI -> Streamlit displays answer, evidence snippets, download option

## 5. Additional Requirements / Notes

- Supports local-only execution (no cloud APIs required)
- Chunk size & top-k configurable for performance
- Handles scanned PDFs with OCR
- Optional Web UI allows demonstration for demo video

6. Deliverables Summary

Deliverable	Description
----- -----	
HLD Document	This PDF/DOCX with system overview and diagrams
Source Code	src/, scripts/, requirements.txt, README.md
Indexed Data	data/processed/ & ChromaDB (data/index/chroma)
Demo Video	Short screen recording showing end-to-end workflow